

# DL0101EN-3-1-Regression-with-Keras-py-v1.0

January 15, 2020

Regression Models with Keras

## 0.1 Introduction

As we discussed in the videos, despite the popularity of more powerful libraries such as PyTorch and TensorFlow, they are not easy to use and have a steep learning curve. So, for people who are just starting to learn deep learning, there is no better library to use other than the Keras library.

Keras is a high-level API for building deep learning models. It has gained favor for its ease of use and syntactic simplicity facilitating fast development. As you will see in this lab and the other labs in this course, building a very complex deep learning network can be achieved with Keras with only few lines of code. You will appreciate Keras even more, once you learn how to build deep models using PyTorch and TensorFlow in the other courses.

So, in this lab, you will learn how to use the Keras library to build a regression model.

## 0.2 Table of Contents

1. Download and Clean Dataset
2. Import Keras
3. Build a Neural Network
4. Train and Test the Network

## 0.3 Download and Clean Dataset

Let's start by importing the pandas and the Numpy libraries.

```
[1]: import pandas as pd
import numpy as np
```

We will be playing around with the same dataset that we used in the videos.

The dataset is about the compressive strength of different samples of concrete based on the volumes of the different ingredients that were used to make them. Ingredients include:

1. Cement
2. Blast Furnace Slag

3. Fly Ash
4. Water
5. Superplasticizer
6. Coarse Aggregate
7. Fine Aggregate

Let's download the data and read it into a pandas dataframe.

```
[2]: concrete_data = pd.read_csv('https://s3-api.us-geo.objectstorage.softlayer.net/
↳cf-courses-data/CognitiveClass/DL0101EN/labs/data/concrete_data.csv')
concrete_data.head()
```

```
[2]:
```

	Cement	Blast Furnace Slag	Fly Ash	Water	Superplasticizer	\
0	540.0	0.0	0.0	162.0	2.5	
1	540.0	0.0	0.0	162.0	2.5	
2	332.5	142.5	0.0	228.0	0.0	
3	332.5	142.5	0.0	228.0	0.0	
4	198.6	132.4	0.0	192.0	0.0	

  

	Coarse Aggregate	Fine Aggregate	Age	Strength
0	1040.0	676.0	28	79.99
1	1055.0	676.0	28	61.89
2	932.0	594.0	270	40.27
3	932.0	594.0	365	41.05
4	978.4	825.5	360	44.30

So the first concrete sample has 540 cubic meter of cement, 0 cubic meter of blast furnace slag, 0 cubic meter of fly ash, 162 cubic meter of water, 2.5 cubic meter of superplasticizer, 1040 cubic meter of coarse aggregate, 676 cubic meter of fine aggregate. Such a concrete mix which is 28 days old, has a compressive strength of 79.99 MPa.

Let's check how many data points we have.

```
[3]: concrete_data.shape
```

```
[3]: (1030, 9)
```

So, there are approximately 1000 samples to train our model on. Because of the few samples, we have to be careful not to overfit the training data.

Let's check the dataset for any missing values.

```
[4]: concrete_data.describe()
```

```
[4]:
```

	Cement	Blast Furnace Slag	Fly Ash	Water	\
count	1030.000000	1030.000000	1030.000000	1030.000000	
mean	281.167864	73.895825	54.188350	181.567282	
std	104.506364	86.279342	63.997004	21.354219	

min	102.000000	0.000000	0.000000	121.800000
25%	192.375000	0.000000	0.000000	164.900000
50%	272.900000	22.000000	0.000000	185.000000
75%	350.000000	142.950000	118.300000	192.000000
max	540.000000	359.400000	200.100000	247.000000

	Superplasticizer	Coarse Aggregate	Fine Aggregate	Age \
count	1030.000000	1030.000000	1030.000000	1030.000000
mean	6.204660	972.918932	773.580485	45.662136
std	5.973841	77.753954	80.175980	63.169912
min	0.000000	801.000000	594.000000	1.000000
25%	0.000000	932.000000	730.950000	7.000000
50%	6.400000	968.000000	779.500000	28.000000
75%	10.200000	1029.400000	824.000000	56.000000
max	32.200000	1145.000000	992.600000	365.000000

	Strength
count	1030.000000
mean	35.817961
std	16.705742
min	2.330000
25%	23.710000
50%	34.445000
75%	46.135000
max	82.600000

```
[5]: concrete_data.isnull().sum()
```

```
[5]: Cement          0
     Blast Furnace Slag  0
     Fly Ash          0
     Water            0
     Superplasticizer  0
     Coarse Aggregate  0
     Fine Aggregate    0
     Age              0
     Strength         0
     dtype: int64
```

The data looks very clean and is ready to be used to build our model.

**Split data into predictors and target** The target variable in this problem is the concrete sample strength. Therefore, our predictors will be all the other columns.

```
[6]: concrete_data_columns = concrete_data.columns
```

```

predictors = concrete_data[concrete_data.columns[concrete_data.columns != 'Strength']] # all columns except Strength
target = concrete_data['Strength'] # Strength column

```

Let's do a quick sanity check of the predictors and the target dataframes.

```
[7]: predictors.head()
```

```

[7]:      Cement  Blast Furnace Slag  Fly Ash  Water  Superplasticizer  \
0      540.0          0.0      0.0  162.0          2.5
1      540.0          0.0      0.0  162.0          2.5
2      332.5        142.5      0.0  228.0          0.0
3      332.5        142.5      0.0  228.0          0.0
4      198.6        132.4      0.0  192.0          0.0

      Coarse Aggregate  Fine Aggregate  Age
0          1040.0          676.0    28
1          1055.0          676.0    28
2           932.0          594.0   270
3           932.0          594.0   365
4           978.4          825.5   360

```

```
[8]: target.head()
```

```

[8]: 0      79.99
1      61.89
2      40.27
3      41.05
4      44.30
Name: Strength, dtype: float64

```

Finally, the last step is to normalize the data by subtracting the mean and dividing by the standard deviation.

```

[9]: predictors_norm = (predictors - predictors.mean()) / predictors.std()
predictors_norm.head()

```

```

[9]:      Cement  Blast Furnace Slag  Fly Ash  Water  Superplasticizer  \
0      2.476712          -0.856472 -0.846733 -0.916319          -0.620147
1      2.476712          -0.856472 -0.846733 -0.916319          -0.620147
2      0.491187          0.795140 -0.846733  2.174405          -1.038638
3      0.491187          0.795140 -0.846733  2.174405          -1.038638
4     -0.790075          0.678079 -0.846733  0.488555          -1.038638

      Coarse Aggregate  Fine Aggregate  Age
0          0.862735          -1.217079 -0.279597
1          1.055651          -1.217079 -0.279597
2         -0.526262          -2.239829  3.551340

```

3	-0.526262	-2.239829	5.055221
4	0.070492	0.647569	4.976069

Let's save the number of predictors to `n_cols` since we will need this number when building our network.

```
[10]: n_cols = predictors_norm.shape[1] # number of predictors
```

## 0.4 Import Keras

Recall from the videos that Keras normally runs on top of a low-level library such as TensorFlow. This means that to be able to use the Keras library, you will have to install TensorFlow first and when you import the Keras library, it will be explicitly displayed what backend was used to install the Keras library. In CC Labs, we used TensorFlow as the backend to install Keras, so it should clearly print that when we import Keras.

Let's go ahead and import the Keras library

```
[11]: import keras
```

Using TensorFlow backend.

As you can see, the TensorFlow backend was used to install the Keras library.

Let's import the rest of the packages from the Keras library that we will need to build our regression model.

```
[12]: from keras.models import Sequential
      from keras.layers import Dense
```

## 0.5 Build a Neural Network

Let's define a function that defines our regression model for us so that we can conveniently call it to create our model.

```
[13]: # define regression model
def regression_model():
    # create model
    model = Sequential()
    model.add(Dense(50, activation='relu', input_shape=(n_cols,)))
    model.add(Dense(50, activation='relu'))
    model.add(Dense(1))

    # compile model
    model.compile(optimizer='adam', loss='mean_squared_error')
    return model
```

The above function create a model that has two hidden layers, each of 50 hidden units.

## 0.6 Train and Test the Network

Let's call the function now to create our model.

```
[14]: # build the model
model = regression_model()
```

Next, we will train and test the model at the same time using the *fit* method. We will leave out 30% of the data for validation and we will train the model for 100 epochs.

```
[15]: # fit the model
model.fit(predictors_norm, target, validation_split=0.3, epochs=100, verbose=2)
```

```
Train on 721 samples, validate on 309 samples
Epoch 1/100
- 0s - loss: 1660.1377 - val_loss: 1143.0640
Epoch 2/100
- 0s - loss: 1523.9059 - val_loss: 1011.1052
Epoch 3/100
- 0s - loss: 1325.1134 - val_loss: 822.5143
Epoch 4/100
- 0s - loss: 1038.2154 - val_loss: 586.4825
Epoch 5/100
- 0s - loss: 698.7738 - val_loss: 362.4619
Epoch 6/100
- 0s - loss: 418.6881 - val_loss: 212.2314
Epoch 7/100
- 0s - loss: 277.2821 - val_loss: 162.6457
Epoch 8/100
- 0s - loss: 231.4284 - val_loss: 159.1913
Epoch 9/100
- 0s - loss: 215.7168 - val_loss: 157.8709
Epoch 10/100
- 0s - loss: 203.6931 - val_loss: 158.1482
Epoch 11/100
- 0s - loss: 195.4887 - val_loss: 156.6341
Epoch 12/100
- 0s - loss: 188.5745 - val_loss: 154.4587
Epoch 13/100
- 0s - loss: 182.9957 - val_loss: 153.1781
Epoch 14/100
- 0s - loss: 178.0344 - val_loss: 149.8669
Epoch 15/100
- 0s - loss: 173.9277 - val_loss: 152.5242
Epoch 16/100
- 0s - loss: 169.5102 - val_loss: 146.3247
Epoch 17/100
- 0s - loss: 165.9091 - val_loss: 148.0968
Epoch 18/100
```

- 0s - loss: 162.7200 - val\_loss: 145.2803  
Epoch 19/100  
- 0s - loss: 159.8348 - val\_loss: 146.9721  
Epoch 20/100  
- 0s - loss: 157.3678 - val\_loss: 143.6480  
Epoch 21/100  
- 0s - loss: 154.4429 - val\_loss: 145.8840  
Epoch 22/100  
- 0s - loss: 152.6101 - val\_loss: 142.0781  
Epoch 23/100  
- 0s - loss: 151.4637 - val\_loss: 141.8954  
Epoch 24/100  
- 0s - loss: 148.4784 - val\_loss: 142.6425  
Epoch 25/100  
- 0s - loss: 146.3285 - val\_loss: 139.9731  
Epoch 26/100  
- 0s - loss: 145.3244 - val\_loss: 140.0204  
Epoch 27/100  
- 0s - loss: 143.2450 - val\_loss: 138.6028  
Epoch 28/100  
- 0s - loss: 141.9220 - val\_loss: 139.6912  
Epoch 29/100  
- 0s - loss: 140.3962 - val\_loss: 140.1303  
Epoch 30/100  
- 0s - loss: 139.2788 - val\_loss: 139.6737  
Epoch 31/100  
- 0s - loss: 138.1343 - val\_loss: 137.6235  
Epoch 32/100  
- 0s - loss: 136.4452 - val\_loss: 139.8918  
Epoch 33/100  
- 0s - loss: 135.5935 - val\_loss: 137.1772  
Epoch 34/100  
- 0s - loss: 133.7596 - val\_loss: 137.5810  
Epoch 35/100  
- 0s - loss: 132.5286 - val\_loss: 140.7972  
Epoch 36/100  
- 0s - loss: 131.1030 - val\_loss: 139.1057  
Epoch 37/100  
- 0s - loss: 130.1083 - val\_loss: 137.1844  
Epoch 38/100  
- 0s - loss: 128.7180 - val\_loss: 137.7035  
Epoch 39/100  
- 0s - loss: 127.7545 - val\_loss: 139.3261  
Epoch 40/100  
- 0s - loss: 126.4812 - val\_loss: 135.9571  
Epoch 41/100  
- 0s - loss: 125.2417 - val\_loss: 137.7308  
Epoch 42/100

- 0s - loss: 123.8980 - val\_loss: 137.2717  
Epoch 43/100  
- 0s - loss: 122.8892 - val\_loss: 136.2065  
Epoch 44/100  
- 0s - loss: 122.1091 - val\_loss: 138.2439  
Epoch 45/100  
- 0s - loss: 120.9098 - val\_loss: 134.9990  
Epoch 46/100  
- 0s - loss: 119.9991 - val\_loss: 137.9615  
Epoch 47/100  
- 0s - loss: 118.3137 - val\_loss: 138.8652  
Epoch 48/100  
- 0s - loss: 117.0378 - val\_loss: 137.3919  
Epoch 49/100  
- 0s - loss: 115.5511 - val\_loss: 134.1835  
Epoch 50/100  
- 0s - loss: 113.8509 - val\_loss: 136.2344  
Epoch 51/100  
- 0s - loss: 112.2659 - val\_loss: 133.5108  
Epoch 52/100  
- 0s - loss: 110.2875 - val\_loss: 134.0008  
Epoch 53/100  
- 0s - loss: 108.5977 - val\_loss: 132.3212  
Epoch 54/100  
- 0s - loss: 106.4394 - val\_loss: 132.3104  
Epoch 55/100  
- 0s - loss: 103.9009 - val\_loss: 130.4516  
Epoch 56/100  
- 0s - loss: 101.3357 - val\_loss: 129.0572  
Epoch 57/100  
- 0s - loss: 98.8990 - val\_loss: 128.3138  
Epoch 58/100  
- 0s - loss: 96.9462 - val\_loss: 127.7728  
Epoch 59/100  
- 0s - loss: 93.2144 - val\_loss: 126.1288  
Epoch 60/100  
- 0s - loss: 90.3830 - val\_loss: 121.7659  
Epoch 61/100  
- 0s - loss: 87.9509 - val\_loss: 121.4203  
Epoch 62/100  
- 0s - loss: 84.7853 - val\_loss: 120.0472  
Epoch 63/100  
- 0s - loss: 81.6367 - val\_loss: 119.1588  
Epoch 64/100  
- 0s - loss: 78.5383 - val\_loss: 114.0035  
Epoch 65/100  
- 0s - loss: 75.4104 - val\_loss: 115.2481  
Epoch 66/100



- 0s - loss: 72.7282 - val\_loss: 108.8653  
Epoch 67/100  
- 0s - loss: 70.4902 - val\_loss: 111.9116  
Epoch 68/100  
- 0s - loss: 68.5002 - val\_loss: 115.8181  
Epoch 69/100  
- 0s - loss: 64.4075 - val\_loss: 107.6068  
Epoch 70/100  
- 0s - loss: 62.4651 - val\_loss: 108.3440  
Epoch 71/100  
- 0s - loss: 60.4697 - val\_loss: 107.7937  
Epoch 72/100  
- 0s - loss: 57.3206 - val\_loss: 102.7240  
Epoch 73/100  
- 0s - loss: 55.7380 - val\_loss: 106.3793  
Epoch 74/100  
- 0s - loss: 53.5651 - val\_loss: 100.9378  
Epoch 75/100  
- 0s - loss: 52.2126 - val\_loss: 102.0156  
Epoch 76/100  
- 0s - loss: 50.3544 - val\_loss: 104.0244  
Epoch 77/100  
- 0s - loss: 48.8139 - val\_loss: 103.0970  
Epoch 78/100  
- 0s - loss: 47.7717 - val\_loss: 109.2335  
Epoch 79/100  
- 0s - loss: 46.8317 - val\_loss: 105.2750  
Epoch 80/100  
- 0s - loss: 45.8606 - val\_loss: 107.8380  
Epoch 81/100  
- 0s - loss: 44.4909 - val\_loss: 114.8449  
Epoch 82/100  
- 0s - loss: 43.1752 - val\_loss: 110.3251  
Epoch 83/100  
- 0s - loss: 42.3568 - val\_loss: 116.2079  
Epoch 84/100  
- 0s - loss: 41.2937 - val\_loss: 109.5871  
Epoch 85/100  
- 0s - loss: 40.4076 - val\_loss: 119.2130  
Epoch 86/100  
- 0s - loss: 39.3967 - val\_loss: 118.4678  
Epoch 87/100  
- 0s - loss: 38.7857 - val\_loss: 118.2535  
Epoch 88/100  
- 0s - loss: 37.8964 - val\_loss: 117.3559  
Epoch 89/100  
- 0s - loss: 38.1799 - val\_loss: 123.1115  
Epoch 90/100

```
- 0s - loss: 37.2709 - val_loss: 123.5457
Epoch 91/100
- 0s - loss: 36.3797 - val_loss: 128.1173
Epoch 92/100
- 0s - loss: 35.5836 - val_loss: 123.2614
Epoch 93/100
- 0s - loss: 34.9860 - val_loss: 128.5807
Epoch 94/100
- 0s - loss: 34.8328 - val_loss: 125.9789
Epoch 95/100
- 0s - loss: 34.3698 - val_loss: 132.5730
Epoch 96/100
- 0s - loss: 34.4769 - val_loss: 128.7075
Epoch 97/100
- 0s - loss: 34.6661 - val_loss: 119.5304
Epoch 98/100
- 0s - loss: 33.3136 - val_loss: 126.3475
Epoch 99/100
- 0s - loss: 33.1318 - val_loss: 126.5016
Epoch 100/100
- 0s - loss: 32.8780 - val_loss: 129.0572
```

[15]: <keras.callbacks.History at 0x7f24b3999860>

You can refer to this [link](#) to learn about other functions that you can use for prediction or evaluation.

Feel free to vary the following and note what impact each change has on the model's performance:

1. Increase or decrease number of neurons in hidden layers
2. Add more hidden layers
3. Increase number of epochs

### 0.6.1 Thank you for completing this lab!

This notebook was created by [Alex Aklson](#). I hope you found this lab interesting and educational. Feel free to contact me if you have any questions!

This notebook is part of a course on **Coursera** called *Introduction to Deep Learning & Neural Networks with Keras*. If you accessed this notebook outside the course, you can take this course online by clicking [here](#).

Copyright © 2019 [IBM Developer Skills Network](#). This notebook and its source code are released under the terms of the [MIT License](#).